

# AWS Elastic Kubernetes Service (EKS) Overview

## 1. Introduction to AWS EKS:

- **AWS EKS** is a managed Kubernetes service that simplifies the process of running Kubernetes on AWS. It automates much of the management overhead, such as setting up the Kubernetes control plane, handling scalability, and ensuring high availability.
- EKS integrates deeply with other AWS services, such as EC2, IAM, and VPC, to provide a secure and scalable environment for Kubernetes workloads.

## 2. Key Features of AWS EKS:

- **Managed Control Plane:** AWS EKS manages the Kubernetes control plane, including the API server, etcd, and other components, across multiple availability zones for high availability.
- **Security and Compliance:** EKS integrates with AWS IAM for role-based access control (RBAC) and supports encrypted data at rest and in transit. EKS clusters are also compliant with various regulatory standards, such as HIPAA, PCI DSS, and more.
- **Auto Scaling:** EKS works with the Kubernetes Cluster Autoscaler and AWS Auto Scaling to automatically adjust the number of nodes in a cluster based on workload demand.
- **Integration with AWS Services:** EKS integrates with various AWS services, including Amazon CloudWatch for monitoring, AWS Load Balancer Controller for load balancing, and AWS Fargate for serverless compute.
- **EKS Add-ons:** AWS EKS allows users to install and manage add-ons like the Amazon VPC CNI, CoreDNS, and kube-proxy.

## 3. Architecture of AWS EKS:

- **Control Plane:** Managed by AWS, it consists of Kubernetes masters distributed across multiple Availability Zones (AZs) for fault tolerance.
- **Worker Nodes:** You can run these on EC2 instances or AWS Fargate. These nodes connect to the control plane via the Kubernetes API and are responsible for running the actual container workloads.
- **Networking:** EKS uses the Amazon VPC Container Network Interface (CNI) plugin for pod networking, enabling pods to receive IP addresses from a VPC subnet and communicate with AWS services.

## 4. Benefits of Using AWS EKS:

- **Simplified Management:** AWS EKS takes care of the Kubernetes control plane, including updates, patching, and scaling, allowing teams to focus more on their applications.
- **Scalability:** EKS can handle the scale of workloads from small applications to enterprise-level, with the ability to scale up and down as needed.
- **High Availability:** EKS runs the control plane across multiple AZs, providing fault tolerance and reducing the risk of downtime.
- **Security:** EKS leverages AWS's robust security features, including IAM, security groups, and encryption, to protect Kubernetes environments.
- **Cost-Effective:** EKS can be more cost-effective than managing Kubernetes clusters manually, as it reduces the operational overhead and complexity.

## 5. Getting Started with AWS EKS:

- **Create an EKS Cluster:** You can create a cluster via the AWS Management Console, AWS CLI, or CloudFormation. This involves specifying the VPC, subnets, and IAM roles.
- **Launch Worker Nodes:** Once the control plane is up, you can launch worker nodes using EC2 instances or Fargate.
- **Deploy Applications:** With the cluster set up, you can deploy Kubernetes workloads using `kubectl` or CI/CD pipelines.
- **Monitor and Manage:** Use Amazon CloudWatch and other monitoring tools to track the performance and health of your EKS cluster.

## 6. Use Cases for AWS EKS:

- **Microservices:** Running microservices-based applications using Kubernetes on AWS.
- **CI/CD Pipelines:** Automating the deployment of applications using Kubernetes as part of CI/CD pipelines.
- **Hybrid Deployments:** Integrating on-premises environments with AWS using Kubernetes.
- **Data Processing:** Running large-scale data processing and analytics workloads.

## 7. Pricing:

- AWS EKS pricing is based on the resources used, including the control plane, worker nodes (EC2 or Fargate), and any additional AWS services consumed. There is a charge for each EKS cluster you create, plus standard charges for other AWS resources (like EC2 instances).

AWS EKS provides a robust, managed Kubernetes environment that integrates seamlessly with the broader AWS ecosystem, making it an ideal choice for organizations looking to deploy containerized applications at scale.

## Setup of a AWS EKS

Setting up an AWS EKS (Elastic Kubernetes Service) cluster involves several steps, including configuring your environment, creating the cluster, launching worker nodes, and deploying applications. Below is a step-by-step guide to setting up AWS EKS:

### 1. Prerequisites

- **AWS Account:** Ensure you have an active AWS account.
- **AWS CLI:** Install and configure the AWS CLI with the necessary IAM credentials.
- **kubectl:** Install `kubectl` to interact with your Kubernetes cluster.
- **eksctl:** Install `eksctl`, a command-line tool to simplify the creation of EKS clusters.  
`# Install eksctl`
- `curl --silent --location "https://github.com/weaveworks/eksctl/releases/latest/download/eksctl_$(uname -s)_amd64.tar.gz" | tar xz -C /tmp`
- `sudo mv /tmp/eksctl /usr/local/bin`

### 2. Create an EKS Cluster

#### Step 1: Configure IAM Roles

- **EKS Cluster Role:** Create an IAM role with the necessary policies (`AmazonEKSClusterPolicy`, `AmazonEKSServicePolicy`).
- **Node Role:** Create an IAM role for the worker nodes with policies like `AmazonEKSWorkerNodePolicy`, `AmazonEC2ContainerRegistryReadOnly`, and `AmazonEKS_CNI_Policy`.

#### Step 2: Create a VPC for the Cluster

- You can either use an existing VPC or create a new one. EKS requires at least two subnets in different Availability Zones for high availability.

```
eksctl create cluster \
```

```
--name my-eks-cluster \  
--region us-west-2 \  
--zones us-west-2a,us-west-2b \  
--nodegroup-name my-eks-nodes \  
--node-type t3.medium \  
--nodes 2 \  
--nodes-min 1 \  
--nodes-max 3 \  
--managed
```

- This command creates a VPC, subnets, an EKS cluster, and a managed node group in the specified region.

### Step 3: Verify the Cluster

- Check the status of your cluster to ensure it was created successfully.

```
eksctl get cluster --name my-eks-cluster
```

- Update your kubeconfig to connect `kubectl` to your EKS cluster:

```
aws eks --region us-west-2 update-kubeconfig --name  
my-eks-cluster
```

## 3. Launch Worker Nodes

- Worker nodes are the EC2 instances that run your containerized applications.
- If not created in the previous step, use the following command:

```
eksctl create nodegroup \  
--cluster my-eks-cluster \  
--region us-west-2 \  
--name my-eks-nodes \  
--node-type t3.medium \  
--nodes 2 \  
--nodes-min 1
```

```
--nodes-max 3
```

## 4. Deploy a Sample Application

- With the cluster and nodes set up, you can deploy a sample application to verify everything is working.

### Step 1: Create a Deployment

- Create a deployment using a sample image:

```
kubectl create deployment hello-world --image=nginx
```

### Step 2: Expose the Deployment

- Expose the deployment to make it accessible:

```
kubectl expose deployment hello-world --type=LoadBalancer  
--port=80
```

- Check the status of the service and note the external IP:

```
kubectl get services
```

## 5. Monitoring and Management

- Use **Amazon CloudWatch** for monitoring logs and metrics.
- You can integrate other monitoring tools like Prometheus and Grafana for more advanced monitoring.

## 6. Clean Up

- When done, clean up resources to avoid unnecessary charges:

```
eksctl delete cluster --name my-eks-cluster
```

This step-by-step guide helps you set up a basic AWS EKS cluster and deploy a sample application. From here, you can explore more advanced configurations and integrations based on your specific use case.